

STEAM REVIEW RADAR

S T A R T



ANDREW WANG

CODY JIANG

DYLAN ZHOU



HOW TO PLAY

THERE ARE 5 PARTS OF
OUR PRESENTATION!

DATASET INTRO

DATA
PREPROCESSING

EDA & FEATURES

SPAM
DETECTION

CONCLUSIONS



PROBLEM MOTIVATION

Steam hosts millions of reviews → critical signal for recommendations and shoppers

But fresh reviews initially have no helpfulness votes, and platform is vulnerable to spam & duplicates.

We want to answer:

- which new reviews will be helpful?
- who are the reviewers behind them?
- how do we detect spam/near-duplicate reviews?


DATASET INTRODUCTION

Dataset Source


- Kaggle: 113,883,717 reviews (~42GB)
- Stored in CSV → Parquet
- Sample extraction sizes:
 - Small ≈ 20k for prototyping
 - Medium ≈ 2M for model training

Fields used

- Review text, helpful votes, funny votes
- Author playtime, ownership, purchase flags
- Timestamps, language, metadata




DATA PREPROCESSING PIPELINE



Processing Steps

1. Chunk-read CSV → downsample to reduce load of next steps
2. Filter for English reviews only
3. Probabilistic sampling → medium (2M) & small (20K) subsets
4. Remove unnecessary columns
5. Convert and store Parquet



Convert to Parquet

Sanity check for our new datasets and filter out some columns, then convert to parquet

```
sml_path = "steam_dataset/sml_sample.csv"
sml = pd.read_csv(sml_path)
sml = sml.drop(['hidden_in_steam_china', 'steam_china_location'], axis='columns', errors='ignore')

print(len(sml))
# sml.head()
# sml.columns
print(sml.dtypes)

sml.to_csv(sml_path, index=False, quoting=csv.QUOTE_NONNUMERIC)

20003
recommendationid           int64
appid                         int64
game                           object
author_steamid                  int64
author_num_games_owned          int64
author_num_reviews                  int64
author_playtime_forever          int64
author_playtime_last_two_weeks    int64
author_playtime_at_review         int64
author_last_played                  int64
language                          object
review                           object
timestamp_created                  int64
timestamp_updated                  int64
voted_up                         int64
votes_up                          int64
votes_funny                       int64
weighted_vote_score                float64
comment_count                      int64
steam_purchase                     int64
received_for_free                   int64
written_during_early_access        int64
dtype: object
```





EDA HIGHLIGHTS & FEATURE ENGINEERING

Key Observations

- Playtime & review count distribution heavily right-skewed
- Purchased reviews dominate; early-access is minority
- Many short reviews ("good", "nice") → low signal density

Features for model

- 📌 Text: TF-IDF (1–2 grams, 3k vocab)
- 📌 Numeric: playtime, reviewer history, funny upvotes...
- 📌 Categorical: language, purchase/early-access flags
- 📌 Target: votes_up (clipped top 1% to reduce long tail)






HELPFULNESS PREDICTION (MODEL RESULTS)


Model

- Poisson Regression pipeline
- Inputs: TF-IDF + metadata + numeric behavior
- Output: predicted helpful upvotes

Performance

- MSE = 10.16
- RMSE = 3.18
- MAE = 1.00
- Spearman $\rho = 0.716 \leftarrow$ strong ranking correlation







REVIEWER ARCHETYPES

Workflow

- Aggregate review → reviewer-level profile
- Standardize numeric metrics
- Clustering using MiniBatch K-Means (k=5)
- Compare behavioral differences



The Five Steam Reviewer Archetypes											
	Cluster	num_reviews	games_owned	playtime_mean	playtime_recent_mean	votes_up_mean	votes_funny_mean	voted_up_ratio	purchase_ratio	free_ratio	ea_ratio
	0	11.13	84.89	11003.98	40.64	1.53	73925.44	0.99	0.99	0.00	0.13
	1	12.61	123.75	14365.44	69.89	1.85	63	0.89	0.33	0.92	0.10
	2	11.35	112.14	13171.32	38.76	1.53	10261.33	0.99	0.05	0.00	0.09
	3	4.85	64.93	231304.35	2264.56	3.03	124677.19	0.90	0.59	0.02	0.07
	4	15.18	127.48	10003.72	35.70	6.97	149207.28	0.05	0.62	0.00	0.13





SPAM DETECTION PIPELINE



WHY WE NEEDED SPAM DETECTION

Steam has millions of user voices – but not all are real signals.

We noticed many reviews were one-liners, templated, or repeated across games.

Those reviews artificially boost helpfulness metrics and pollute training.

If we train on them directly → our helpfulness model learns wrong behaviors.

Our Questions

- How much of the review ecosystem is copypasta / near-duplicate?
- Can we automatically flag suspicious reviews at scale?
- What if we remove or down-weight them – do predictions improve?

Goal

- Build a lightweight, scalable spam radar that runs on millions of reviews
- Without deep neural embeddings or heavy compute

=====

Pair: 899 vs 12454

[Review A]:
good

[Review B]:
good


=====

Pair: 1829 vs 18023

[Review A]:
great game

[Review B]:
great game

	review_clean	word_count	unique_ratio	avg_word_len
0	thumbsup	1	1.000000	8.000000
1	there s like no guidance on what to do i got a...	46	0.847826	3.239130
2	2d minecraft never felt so good	6	1.000000	4.333333
3	the battlepass is heresy	4	1.000000	5.250000
4	oooooo	1	1.000000	6.000000



SPAM DETECTION PIPELINE

HOW WE DETECT SPAM

We started small – cleaned 20k reviews first, so we could iterate fast.

We processed raw text into normalized tokens to compare reviews fairly.

1. Clean Text

- lowercase, remove symbols, collapse whitespace
- Example: ":thumbsup:" → "thumsup"

2. Generate 5-gram shingles

- Turn each review into overlapping character chunks
- "great game" → ["great", "reat ", "eat g", "at ga", "t gam", " game"]

3. MinHash Signatures

- Hash shingles into compact fingerprints
- Makes large-scale similarity search efficient

4. LSH (Locality Sensitive Hashing)

- Bucket reviews with similar fingerprints
- Near-duplicate threshold ≈ 0.8 Jaccard

5. Output: Pairs of suspiciously similar reviews

Evidence of success

➤ On 19,770 reviews, 87,267 duplicate pairs found

```
# Build MinHash signatures and LSH index
num_perm = 128

def build_minhash(shingles):
    """
    Build a MinHash signature for a list of shingles.
    """
    m = MinHash(num_perm=num_perm)
    for s in shingles:
        if s: # skip empty strings
            m.update(s.encode("utf8"))
    return m

df = df.reset_index(drop=True)

minhashes = {}
for idx, shingles in df["shingles_5gram"].items():
    m = build_minhash(shingles)
    minhashes[idx] = m

print("Number of MinHash signatures built:", len(minhashes))

# Build MinHash LSH index
threshold = 0.8

lsh = MinHashLSH(threshold=threshold, num_perm=num_perm)

for idx, m in minhashes.items():
    lsh.insert(str(idx), m)

print("LSH index built.")
```

LSH index built.

```
# Query the LSH index to find near-duplicate pairs

pairs = set()

for idx, m in minhashes.items():
    candidates = lsh.query(m)
    i = idx
    for c in candidates:
        j = int(c)
        if j == i:
            continue
        a, b = sorted((i, j))
        pairs.add((a, b))

print("Unique near-duplicate pairs found:", len(pairs))

sample_pairs = list(pairs)[:10]
for (a, b) in sample_pairs:
    print("=" * 80)
    print(f"Pair: {a} vs {b}")
    print("\n[Review A]:")
    print(df.loc[a, "review_clean"])
    print("\n[Review B]:")
    print(df.loc[b, "review_clean"])

Unique near-duplicate pairs found: 87267
```

SPAM DETECTION PIPELINE

WHAT WE FOUND + WHY IT MATTERS

Once we flagged duplicates, we measured how different they actually were.

And the results confirmed what we suspected – spam reviews behave differently.

Key Findings

- 14.74% of reviews were near-duplicate
- Style anomaly (IsolationForest) caught ~2% unusual-writing reviews
- Spam flagged reviews had much lower helpful votes:

Group Avg votes_up

Normal reviews 3.09

Near-duplicate spam 0.54

Style anomalies 17.8* (rare but extreme cases)

Why it matters




- Spam reviews drag down regression training
- We can drop them, down-weight them, or add flags as model features
- Even baseline Poisson regression improved when spam was removed

Future

- Use transformer embeddings for semantic copies
- Live spam radar could run in stream (Kafka → Spark Structured Streaming)

"By removing spam, we clean the signal.

By boosting real reviews, we support authentic player voices."




CONCLUSION & TAKEAWAYS



Achievements

- 71.6% rank correlation helpfulness predictor
- Reviewer archetypes reveal platform behavioral patterns
- 14.7% duplicate detection improves data quality

Future Improvements

- Train XGBoost Poisson / GBDT
- Incorporate semantic embeddings (BERT)
- Deploy streaming real-time scoring agent





THANK YOU!
QUESTIONS WELCOME

